

Figure 10 - Efficiency: model size and CPU latency

